

Dynamic Allocation




Dynamic Storage

- Example: We need space to allow the user to enter some data, but how much?
 - 80 bytes? 132? 1? 256? 10000?
- Two Friends from the C library: `#include <stdlib.h>`
 - `malloc()`
 - `free()`
- These friends create and manage the heap for us, allowing us to grab storage at run time

```
struct r {  
    char name[40];  
    double gpa;  
    struct r *next;  
};  
struct r *rp;  
rp = malloc(sizeof(struct r));  
if (rp == NULL){  
    /* Handle Error! */  
}
```

➤ Options for handling error

- Abort
- Ask again
- Save user data
- Ask for less
- Free up something



Don't Do This!

```
rp = malloc(sizeof struct r);
```

```
// Code foolishly inserted here!
```

```
if(rp == NULL)
```

➤ Be afraid.

➤ Using **rp** before you check it is a crash waiting to happen.

Idiomatic But Safe

```
if((rp = malloc(sizeof struct r)) == NULL) {  
    /* Handle Error Here */  
}
```

- This is the way it's done in real life.
- Don't mix up the = and == operators!
- Alternate syntax (*because NULL==0*):

```
if( !(rp = malloc(sizeof struct r)) ){  
    /* Handle Error Here */  
}
```

Escape From Type-Checking

- malloc() returns a pointer to at least as many bytes as we requested.
- But what is the type of this pointer and how do we use it to store items of a different type?
- malloc() is declared as “void *malloc(unsigned long)”
- C uses the idiom “pointer to void” for a generic pointer
- To be safe, you should cast this pointer into the correct type so that type-checking can work for you again!

```
int *ip;  
ip = (int *)malloc(sizeof int);
```

- Otherwise, the compiler will silently cast your “void *” pointer into any other kind of pointer without checking

Done With a Chunk of Storage

- When you're done with a chunk of storage, you use **free()** to make it available for reuse.
- Remember, C doesn't do garbage collection
- From our previous example

```
free(rp) ;
```

returns the memory back to the heap for re-use by someone else

- You **must not** use the value in `rp` after the call to **free()**, nor may you dereference the memory it points to!
- There's no guarantee what's at `*rp` after you call `free()` – assume it is garbage data!
- With modern C libraries, `free(NULL)` does nothing but isn't an error

- From our previous example

```
free (rp) ;
```

- The variable `rp` still exists.
 - It is a pointer to struct `r`
- But what it points to is now garbage data.
 - We should never dereference `rp` again: `*rp`
- We can, however, assign a new value to `rp`
 - That is okay – make it point somewhere else
- The compiler will NOT help you with this.
 - Mistakes will cause run-time errors

Other memory allocation functions

➤ `void *malloc(size_t n);`

- Allocates (at least) `n` bytes of memory on the heap, returns a pointer to it
- Assume memory contains garbage values

➤ `void *calloc(size_t num, size_t size);`

- Allocates (at least) `num*size` bytes of memory on the heap, returns a pointer to it
- Memory will be zero'ed out.

➤ `void *realloc(void *ptr, size_t n);`

- Reallocates (at least) `n` bytes of memory on the heap, returns a pointer to it
- Copies the data starting at `ptr` that was previously allocated
- Often used to expand the memory size for an existing object on the heap

➤ <https://www.youtube.com/watch?v=5VnDaHBi8dM>

Handling Persistent Data



```
char *foo(void)
{
    static char ca[10];
    return ca;
}
```

- Anyone calling this function now has access to ca in this block. Could be dangerous. Why?
- Note that this approach is not dynamic

Example

```
char *strFromUnsigned(unsigned u)
{
    static char strDigits[] = "?????" ;

    char *pch;

    pch = &strDigits[5];

    do
        *--pch = (u % 10) + '0';

    while((u /= 10) > 0);

    return pch;
}
```

```
strHighScore =  
    strFromUnsigned(HighScore) ;
```

.

.

.

```
strThisScore =  
    strFromUnsigned(ThisScore) ;
```

```
char *foo(void)
{
    char ca[10];
    return ca;
}
```

➤ Since ca was allocated on stack during function call pointer returned is now pointing to who-knows-what

➤ **Bad**

```
char *foo(void)
{
    char *ca = malloc(...);
    /* error checking but no free */
    return ca;
}
```

- This actually works, but the caller needs to know that they're responsible for the free()

Memory Leaks

- Memory leaks occur when the programmer loses track of memory allocated by malloc or other functions that call malloc

```
void foo(void)
{
    char *ca = malloc(...);
    /* no free */
    return;
}
```

➤ Bad

Memory Management

- Some functions that call malloc
 - calloc
 - strdup
 - regcmp
 - others...
- C doesn't do automatic memory management for efficiency reasons
 - If you want to manage memory...do it yourself!


```
void *calloc(size_t num, size_t size);
```

- Call malloc() to find space for **num** new allocations
- Initialize the space to zero (0)
- Not much to discuss...


```
ptr = realloc(ptr, num_bytes);
```

- What it does (conceptually)
 - Find space for new allocation
 - Copy original data into new space
 - Free old space
 - Return pointer to new space

realloc()



realloc: What might happen



➤ Realloc may return

- same pointer
- different pointer
- NULL

➤ Is this a good idea?

```
cp = realloc(cp, n) ;
```

1. Yes
2. No
3. Sometimes

➤ Is this a good idea?

```
cp = realloc(cp, n);
```

➤ No!

➤ If realloc returns NULL cp is lost

➤ Memory Leak!

How to do it properly

```
void *tmp;  
  
if((tmp = realloc(cp,...)) == NULL)  
{  
    /* realloc error */  
}  
  
else  
{  
    cp = tmp;  
    free(tmp);  
}
```



NO!

Additional Edge Cases

➤ `realloc(NULL, n) ≡ malloc(n);`

➤ `realloc(cp, 0) ≡ free(cp);` // only on some compilers

➤ These can be used to make `realloc` work in a single loop design to build a dynamic structure such as a linked list.

Example

```
int size = 0;      /* Size of "array" */
int *ip = NULL;    /* Pointer to "array" */
int *temp;
int i;
char buffer[80];
while(fgets(buffer, 80, stdin) != NULL) {
    size++;
    if((temp = realloc(ip, size*sizeof(*temp))) == NULL){
        fprintf(stderr, "Realloc failure\n");
        exit(EXIT_FAILURE);
    }
    ip = temp;
    ip[size-1] = strtol(buffer, NULL, 10);
}
```

Dynamic Allocation: What can go wrong?

- Allocate a block and lose it by losing the value of the pointer
- Allocate a block of memory and use the contents without initialization
- Read or write beyond the boundaries of the block
- Free a block but continue to use the contents
- Call realloc to expand a block of memory and then – once moved – keep using the old address
- FAIL TO NOTICE ERROR RETURNS

Questions?




Question

Is there anything wrong with the following code that frees all the nodes in a linked list?

```
for (struct node *p = Head; p != NULL; p = p->next)
    free(p);
```

```
struct node *t;
for (struct node *p = Head; p != NULL; p = t) {
    t = p->next
    free(p);
}
```

- A. The head of the list doesn't get freed
- B. The code is satisfactory as it stands
- C. The reinitialization uses memory that has been freed 
- D. The code uses free() multiple times on the same pointer

Malloc Implementation



➤ K&R Malloc Implementation

- Headers
- Heap Layout
- Allocating the Correct Amount of Memory
- `malloc()` : Getting the memory for work
- `free()` : Recycling the memory when done
- `morecore()` : OS requests for real memory

Memory Management


- The K&R implementations of malloc() and free() are one of many ways to implement these.
- We've refactored the code a bit. It was written for compactness and efficiency in the 1970s C environment; things have changed and we can afford to be a little bit more verbose for clarity.
- It's not necessarily the most efficient, and certainly not the only way:
 - <http://codinghighway.com/2013/07/13/the-magic-of-wrappers/>
 - http://www.inf.udec.cl/~leo/Malloc_tutorial.pdf
 - <https://danluu.com/malloc-tutorial/>

How K&R Malloc Works

- Linked list to track free memory

- Located in the Heap

- *Circular* Linked List



- This is the ***Free List***

- Contains memory available to be malloc'ed

- Once memory is malloc'ed, we don't track it until it is free'd

How K&R Malloc Works

➤ `void *malloc(size_t n);`

➤ Delete a node from the free list

➤ Return a pointer to that node to the program

➤ `void free(void *ptr)`

➤ Insert the node ptr into the free list

➤ Note: We've renamed them `_malloc()` and `_free()` in our code to keep from conflicting with the `malloc()` and `free()` from the C library

➤ There are a number of C library functions that use `malloc()` and `free()`; we don't want override the library versions until our versions work properly!

Question

In the K&R implementation of malloc(), the data structure in the heap is a:

- A. Binary Tree
- B. Hash Table
- C. Circular Linked List
- D. Array of Structs



Question

In the K&R implementation of the heap management, free() does the following:


- A. Deletes a node from the free list
- B. Inserts a node in the free list
- C. Zeroes out the bytes on the heap
- D. Reduces the size of the heap



How K&R Malloc Works

- Each node on the free list is available memory
 - Everything on the free list is in heap memory areas that malloc has requested from the system
 - Other functions can request heap space from the system, but that space will never show up on malloc's free list
- Consider this:
 - We want each block in the free list to be as large as possible
 - So it is more likely to have enough bytes to satisfy a malloc() call
 - Therefore when we free a block adjacent to our other memory on the heap, we should merge that block into the adjacent block.

A Node in the Free List



A block returned by malloc

How K&R Malloc Works

➤ Design choices

- Linked List is ordered by memory address (in the heap)
- We could also order the linked list from largest to smallest size
- There are many other ways we could implement malloc()


➤ Size in K&R malloc is number of *units*

- Not number of bytes
- A block is sizeof(Header) rounded up to the strictest alignment boundary
- A block is 16 bytes in this implementation

Section of K&R Code

```
struct header {                /* block header */
    struct header *next;       /* next block if on free list */
    unsigned size;             /* size of this block in Headersize units */
};
typedef struct header Header;

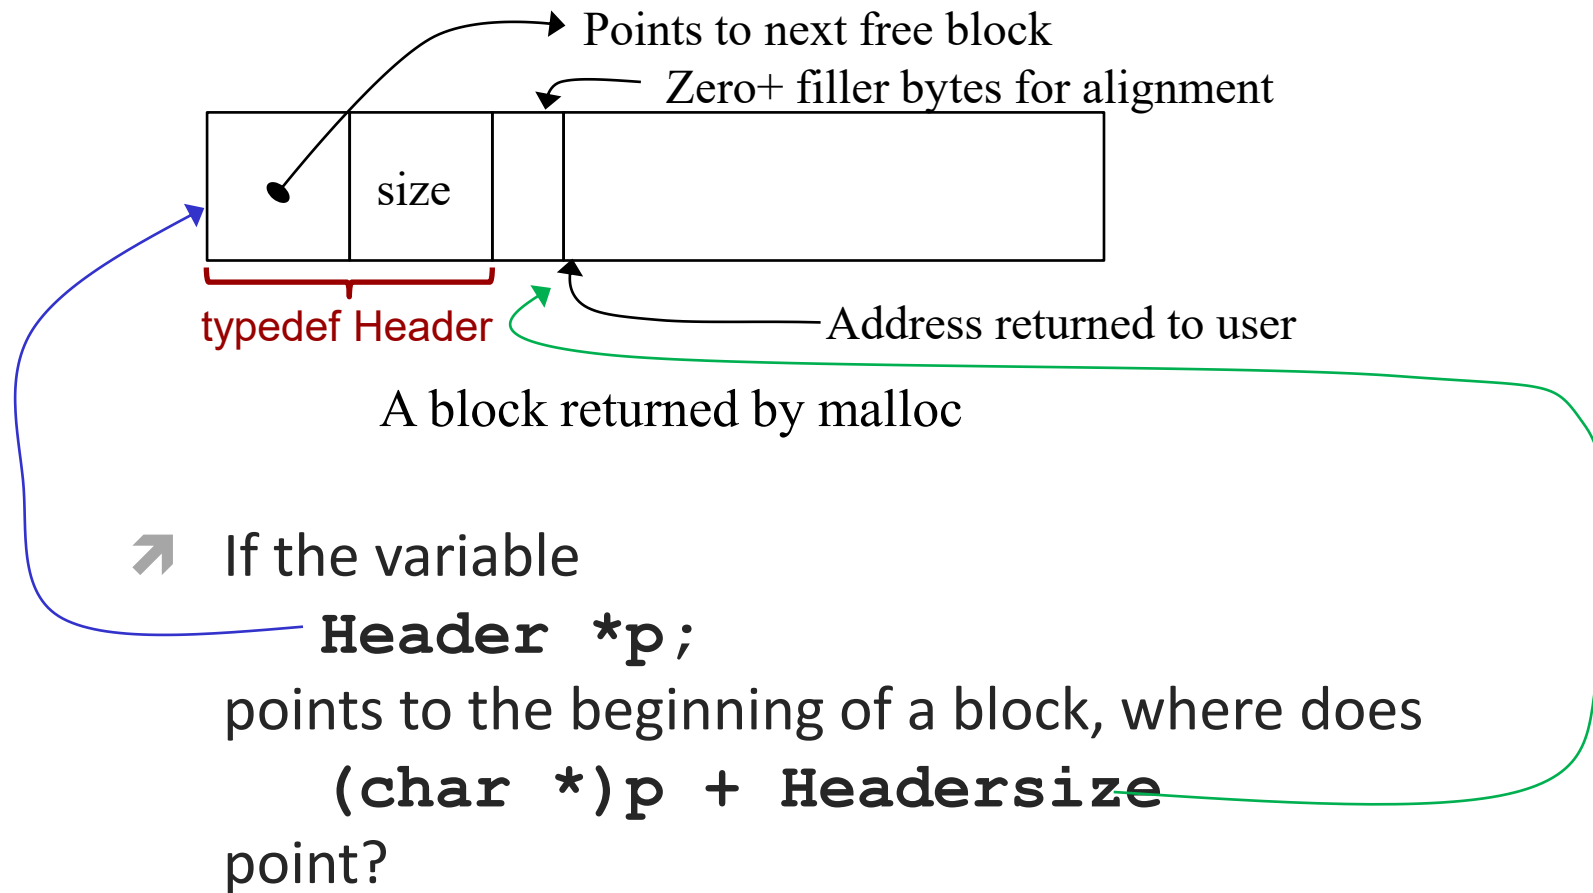
#define ALIGNBOUNDARY 16      /* align memory on 16 byte boundaries */
static int Headersize        /* rounded up; use in place of sizeof(Header) */
    = ((sizeof(Header) - 1) / ALIGNBOUNDARY + 1) * ALIGNBOUNDARY;
```



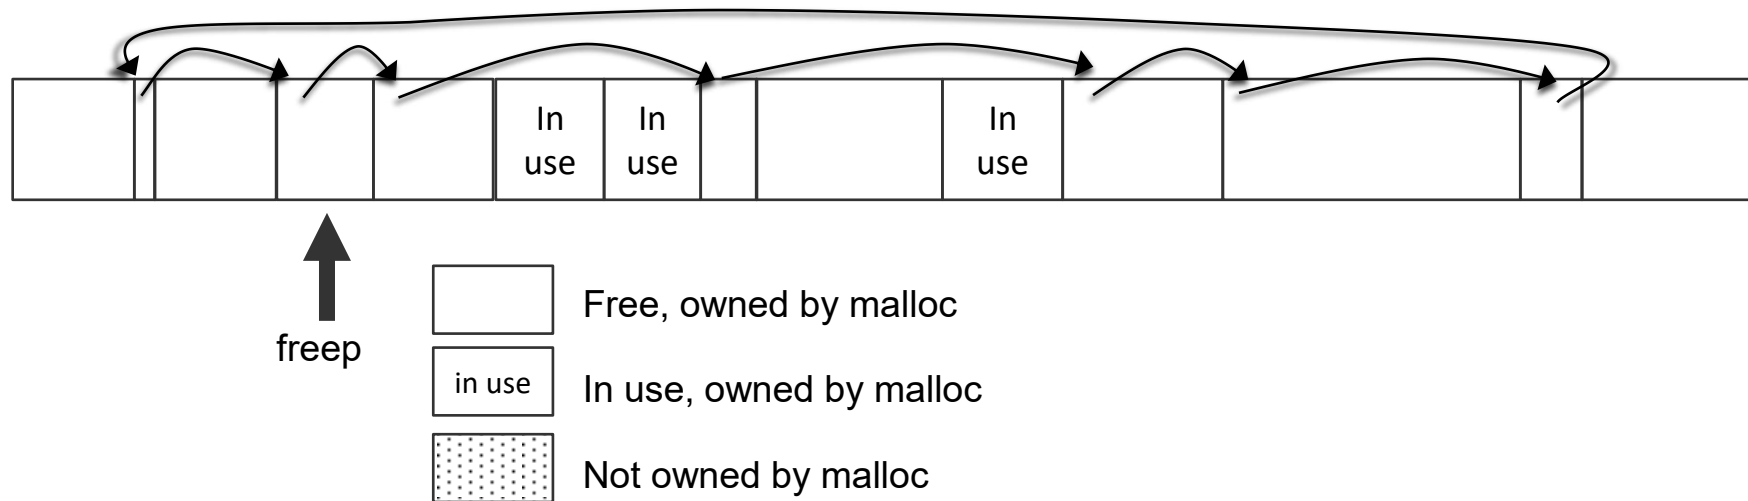
A block returned by malloc

sizeof(Header)	Headersize
8	16
12	16
16	16
20	32

Returned by Malloc



Heap Layout



➤ K&R ch. 8

➤ One of those free blocks has user data size 0 and isn't in the heap, so it is never handed over by malloc().

Section of K&R Code

```
static Header base;                // empty list to get started with
static Header *freep = NULL;       // start of free list

/* malloc: general-purpose storage allocator */
void *malloc(unsigned nbytes)
{
    Header *p, *prevp;
    unsigned long nunits;


    /* Allocate memory in Headersize units, so round up request */
    nunits = (nbytes + Headersize - 1) / Headersize + 1;
```

- The key here is alignment with the Header's size
- You need (just) enough bytes to cover the program's request and the header information
- Suppose that the Header consists of 16 bytes [Headersize = 16]
- Take a look at different requests for sizes as represented by nbytes

nbytes	7	15	16	17	32	33
nunits	2	2	2	3	3	4


Section of malloc() Code

```
if (freep == NULL) {  
    freep = &base;  
    base.size = 0;  
    base.next = &base;  
    if (morecore(nunits) == NULL)  
        return NULL;    /* none left; give up */  
}
```




Section of malloc() Code

```
for (prevp = freep, p = prevp->next; p->size < nunits;  
    prevp = p, p = p->next) {  
    if (p == freep)  
        if ((p = morecore(nunits)) == NULL)  
            return NULL;    /* none left; give up */  
}
```



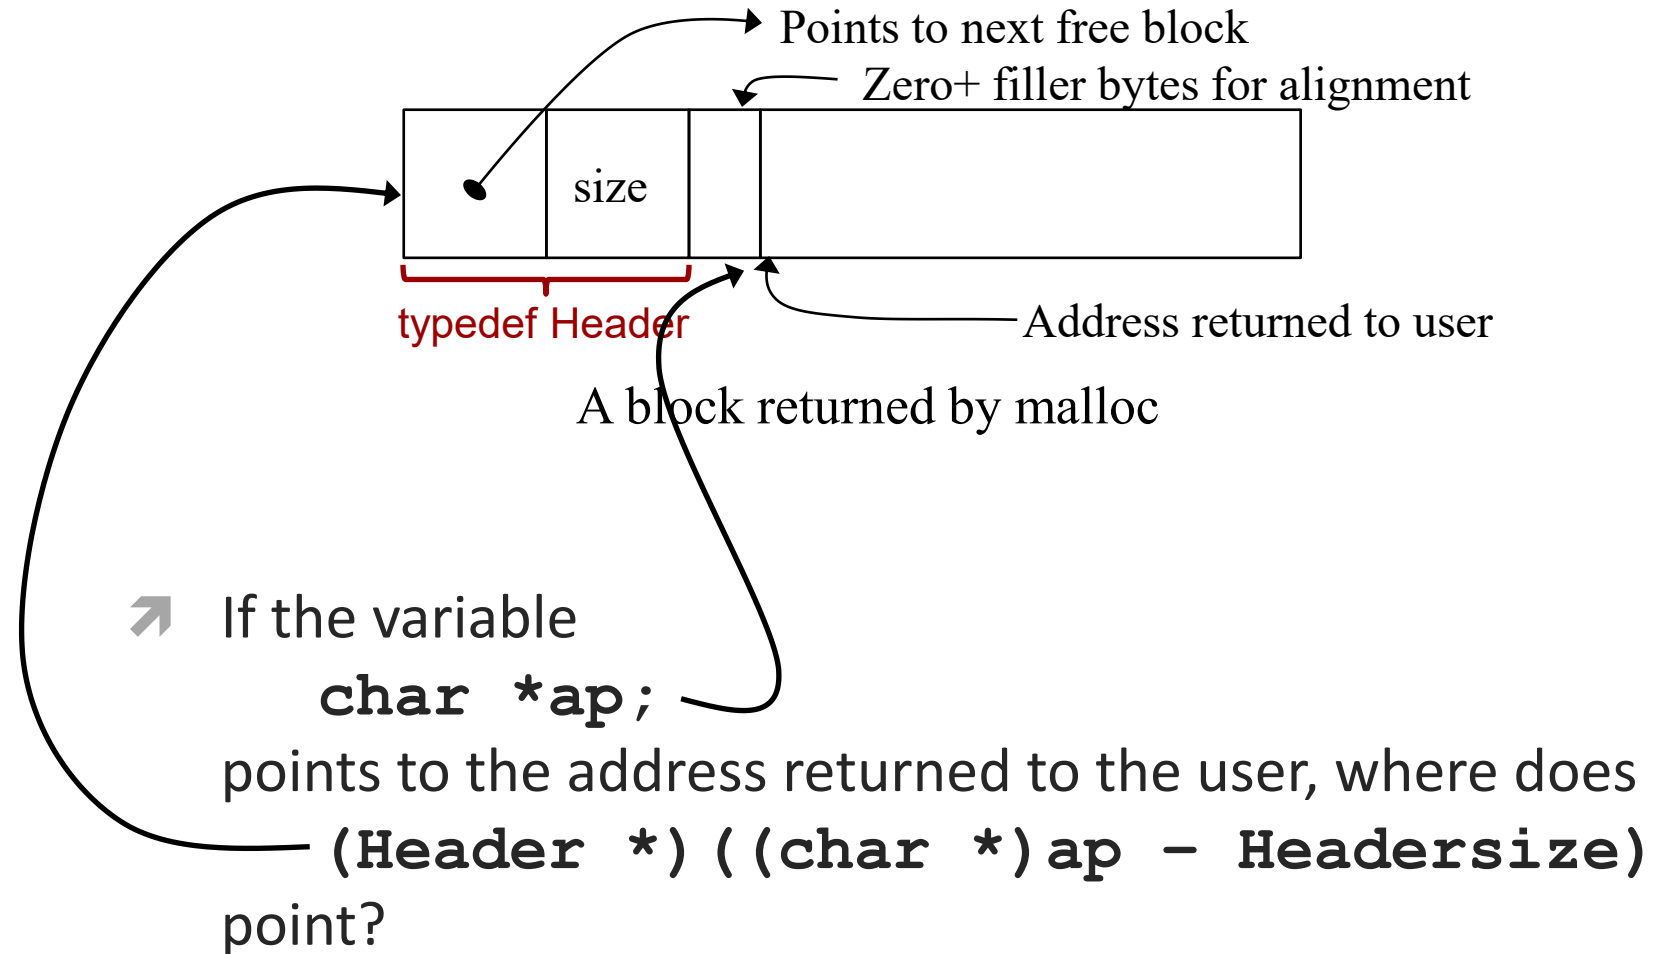
Section of malloc() Code

```
if (p->size == nunits)
    prevp->next = p->next;
else {
    p->size -= nunits;
    p = (char *)p + p->size * Headersize;
    p->size = nunits;
}
freep = prevp;
return (char *)p + Headersize;
```



Let nunits = 7, and follow how the free block is identified and returned

Pointer Sent to free()




Section of free() Code

```
void _free(void *ap)
{
    Header *bp, *p;

    bp = (Header *) (ap - Headersize);


    /* Find p such that the freed block is between p and p->next */
    for (p = freep; !(p < bp && bp < p->next); p = p->next)
        /* is *bp at either end of the arena? */
        if (p >= p->next && (bp > p || bp < p->next))
            break;
```



Free(): Below and Above In Use

```
/* Look to see if we're adjacent to the block after the freed block */
if ((char *)bp + bp->size * Headersize == (char *)p->next) {
    bp->size += p->next->size;
    bp->next = p->next->next;
} else
    bp->next = p->next;
```

```
/* Look to see if we're adjacent to the block before the freed block */
if ((char *)p + p->size * Headersize == (char *)bp) {
    /* add the freed block to the block at p */
    p->size += bp->size;
    p->next = bp->next;
} else
    p->next = bp;
freep = p;
```



Free(): Below In Use

```
/* Look to see if we're adjacent to the block after the freed block */  
if ((char *)bp + bp->size * Headersize == p->next) {
```

```
    bp->size += p->next->size;  
    bp->next = p->next->next;
```

```
} else  
    bp->next = p->next;
```

```
/* Look to see if we're adjacent to the block before the freed block */  
if ((char *)p + p->size * Headersize == bp) {
```

```
    /* add the freed block to the block at p */
```


```
    p->size += bp->size;
```

```
    p->next = bp->next;
```

```
} else
```

```
    p->next = bp;
```


```
    freep = p;
```



Free(): Above In Use

```
/* Look to see if we're adjacent to the block after the freed block */  
if ((char *)bp + bp->size * Headersize == p->next) {  
    bp->size += p->next->size;  
    bp->next = p->next->next;  
} else  
    bp->next = p->next;
```

```
/* Look to see if we're adjacent to the block before the freed block */  
if ((char *)p + p->size * Headersize == bp) {  
    /* add the freed block to the block at p */  
    p->size += bp->size;  
    p->next = bp->next;  
} else  
    p->next = bp;  
freep = p;
```




Free(): Above and Below Free

```
/* Look to see if we're adjacent to the block after the freed block */  
if ((char *)bp + bp->size * Headersize == p->next) {
```

```
    bp->size += p->next->size;  
    bp->next = p->next->next;  
} else  
    bp->next = p->next;
```

```
/* Look to see if we're adjacent to the block before the freed block */  
if ((char *)p + p->size * Headersize == bp) {
```

```
    /* add the freed block to the block at p */  
    p->size += bp->size;  
    p->next = bp->next;  
} else  
    p->next = bp;  
freep = p;
```

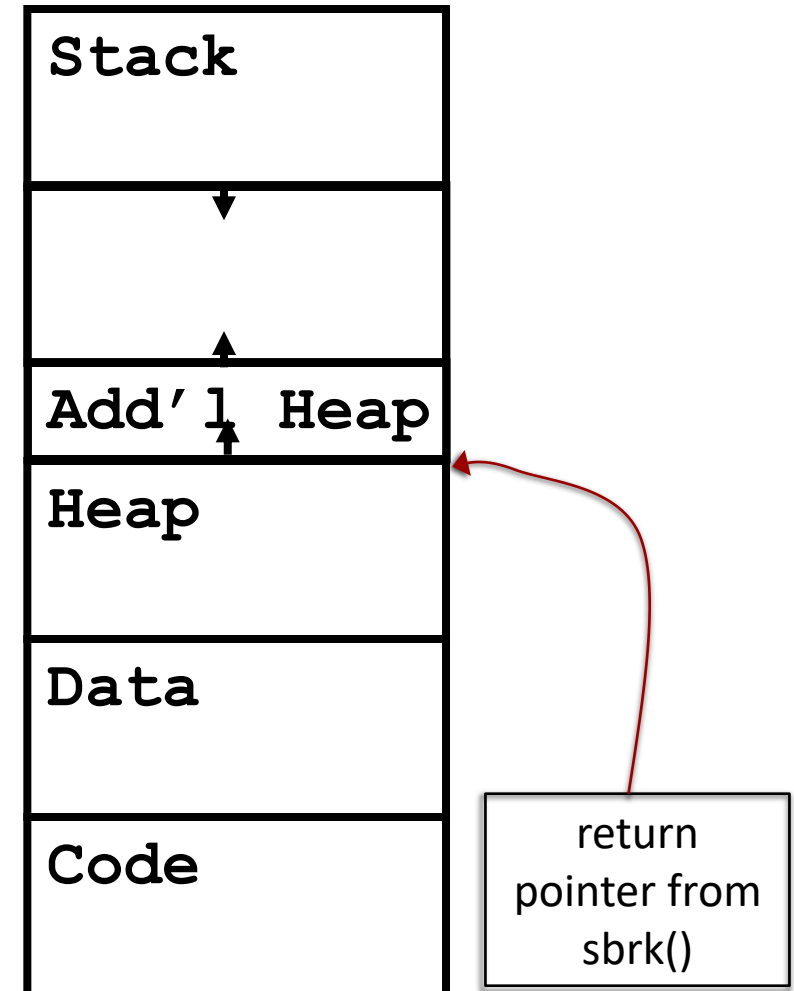


Section of morecore() Code 2

```
/* Don't ask for any less than NALLOC */
if (nu < NALLOC)
    nu = NALLOC;
cp = sbrk(nu * Headersize);
if (cp == (char *) -1) /* no space */
    return NULL;

/* Take the new block, add a header, and free() it */
up = (Header *) cp;
up->size = nu;
_free((char *)up + Headersize);
return freep;
```

- sbrk() is a Linux system call to obtain more memory.
- It works at the end of the data segment/heap (a.k.a. “the break”) to acquire at least n bytes more space from the OS
- It returns an aligned pointer to the new space or -1 if no space is available
- The -1 is an odd historical artifact



Question

After the execution of this code,

```
double *mp = &d[3];  
char *xp = (char *)mp + 16;
```

to what memory address does *xp* point?

- A. 0x10 bytes past the address of *d[3]*
- B. 10 bytes past the beginning of *d*
- C. 10 bytes before the beginning of *d*
- D. 16 bytes before the beginning of *d*



Memory Management

- What are the equivalents of malloc and free in Java? Python?
- How can we “automatically” collect garbage?
 - Mark & Sweep Techniques
 - Reference Counting Techniques
 - And many, many others...