

The ARM ISA for LC-3 Programmers

Copyright © Thomas M. Conte
with sources from ARM, Ltd.
Revised William D. Leahy Jr.

Today's number is 72,000

The claim we made was...

LC-3 is easy to learn

We will teach you LC-3, and *then it will be “easy” to learn another, real ISA*

Time to show you we were right!

A thought: *A new ISA is like a new car. The pedals are there, there's a steering wheel, etc. But the headlight control, high beams, wipers, etc, are all in a slightly different place. Still, you can drive it without too much trouble, and not (hopefully) crash into something!*

REVIEW: ISA= Instruction Set Architecture

ISA = The binary “language” used to talk to a processor. All of the *programmer-visible* components and operations of the computer

- **memory organization**
 - address space -- how many locations can be addressed?
 - addressability -- how many bits per location?
- **register set**
 - how many? what size? how are they used?
- **instruction set**
 - Operation codes)
 - data types (2's complement, floating-point, etc)
 - addressing modes.

Some ARMisms

ARM was a small company in the UK. It was sold to SoftBank in 2016 and is being bought now by NVIDIA.

Technically, they make nothing – they sell the rights to use ARM-compliant processors in your design!

Originally called Acorn Computer Company

The Acorn was a small, cheap computer that used the Mostek 6502

Acorn decided to make their own processor: the Acorn RISC Machine (A.R.M.)

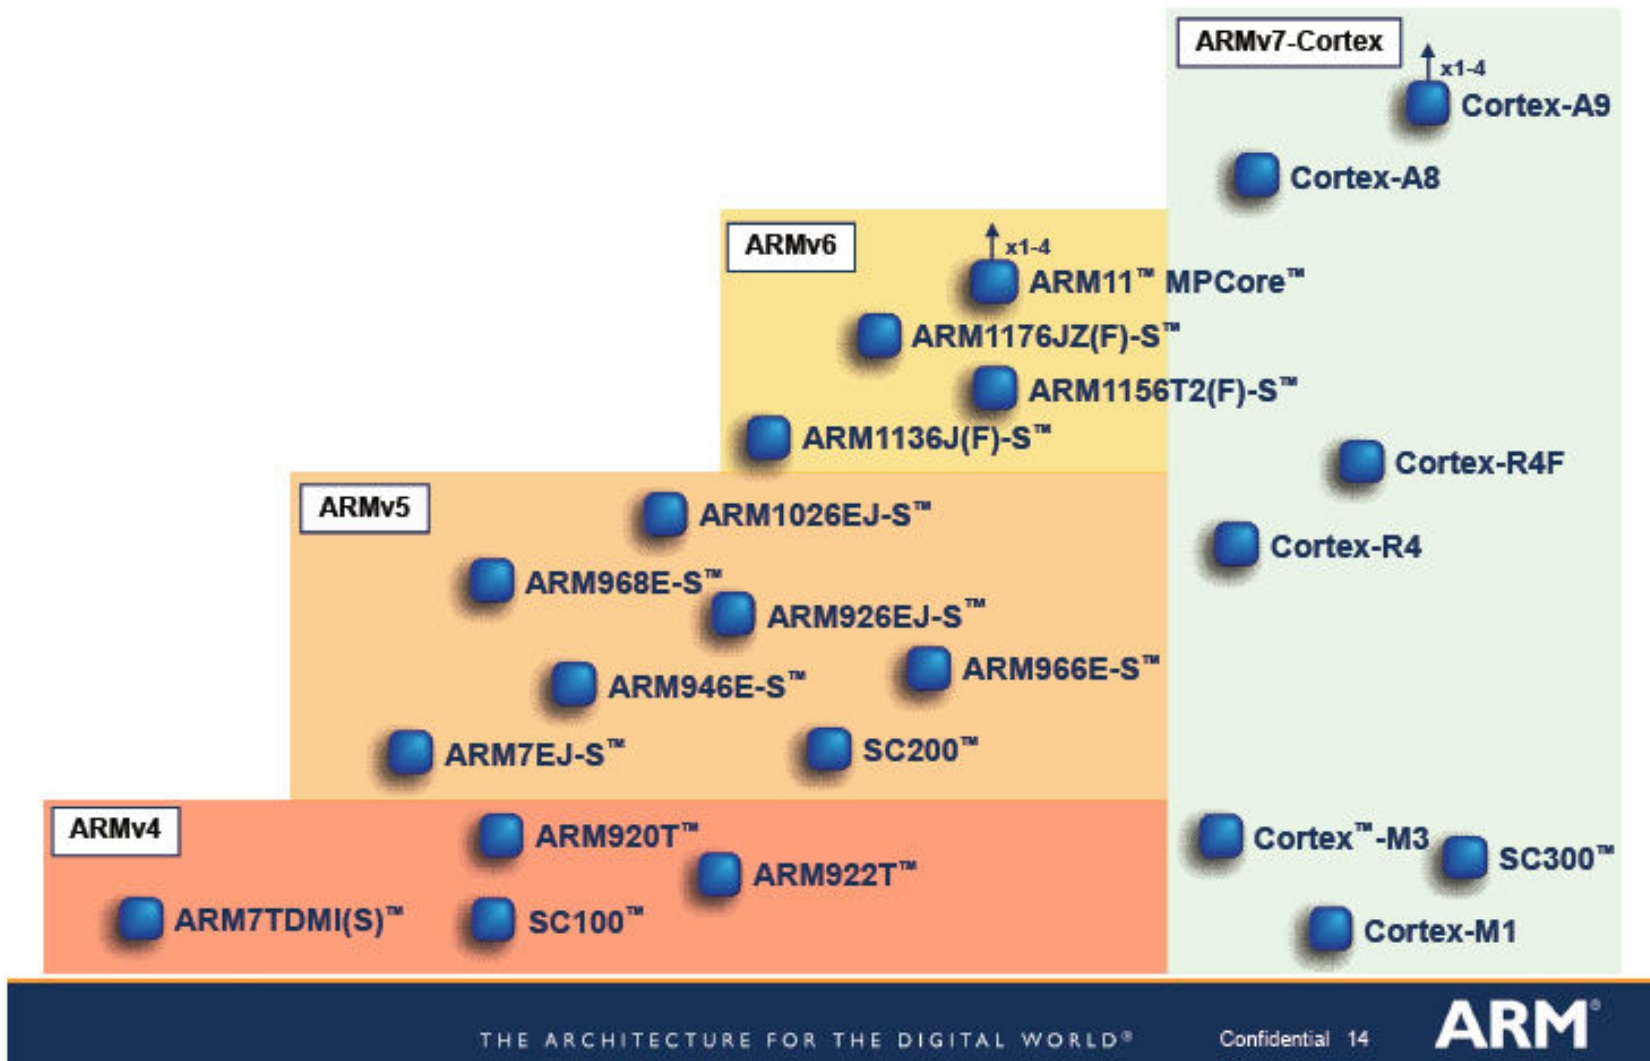
Since that time, A has become “Advanced” instead of “Acorn,” and company changed its name to “ARM”

ARM has had multiple *versions*

The *version* != the processor number

So ARM-7 does NOT use ARM-version7's ISA (it uses ARM version 4)

Architecture Versions



What's the same as LC-3?

ARM and LC-3 are both “load/store” architectures

This means that values get into registers using a LOAD instruction and go to memory using a STORE instruction

ARM and LC-3 both have fixed-format instruction encodings

All instructions in LC-3 are 16b long. ARM has two ISAs, in one, they are 32b long (what I'll present), and they are 16b in the other

Register-register addressing

In ARM as in LC-3, many (most) instructions take two source registers and store their results in a destination register

Address space, addressability

LC-3 Address Space is 2^{16} words, where a word is 16b

all locations are word addressable

the address size is 16 bits

ARM Address Space is 2^{32} bytes

all locations are byte-addressable

the address size is 32 bits

Processor Modes*

User	Normal execution mode
FIQ	High priority (fast) interrupt request
IRQ	General purpose interrupts
Supervisor	Protected mode for O/S
Abort	Used for memory access violations
Undefined	Used to handle undefined instructions
System	Runs privileged O/S tasks

LC-3 just has User and Supervisor modes

***Modes are entered via interrupts or software instructions**

The ALU and the registers

ARM has 16 registers, R0-R15

UNLIKE LC-3, *the PC is a register that the programmer can access directly* (PC = R15 in ARM)

All registers are 32b wide

On older ARM ISAs (versions 1-5), storing to PC has “unpredictable results”

Other special purpose registers:

R11 (“FP”) is the frame pointer (similar function to R5 in LC-3)

R14 (“LR”) is the link register (same function as R7 in LC-3, it holds the return address for a subroutine call)

R13 (“SP”) is the stack pointer (same function as R6 in LC-3, it holds the address of the top of stack)

The ARM Register Set: Register Banking

Current Visible Registers

User Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr

Banked out Registers

FIQ	IRQ	SVC	Undef	Abort
r8				
r9				
r10				
r11				
r12				
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
spsr	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

FIQ Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr
spsr

User

r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)

Banked out Registers

IRQ

SVC

Undef

Abort

r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

User Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr

Banked out Registers

FIQ	IRQ	SVC	Undef	Abort
r8				
r9				
r10				
r11				
r12				
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
spsr	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

IRQ Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr
spsr

Banked out Registers

User FIQ SVC Undef Abort

	r8			
	r9			
	r10			
	r11			
	r12			
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

User Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr

Banked out Registers

FIQ	IRQ	SVC	Undef	Abort
r8				
r9				
r10				
r11				
r12				
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
spsr	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

SVC Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr
spsr

Banked out Registers

User

FIQ

IRQ

Undef

Abort

r13 (sp)
r14 (lr)

r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)

r13 (sp)
r14 (lr)

r13 (sp)
r14 (lr)

r13 (sp)
r14 (lr)

spsr

spsr

spsr

spsr

The ARM Register Set

Current Visible Registers

User Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr

Banked out Registers

FIQ	IRQ	SVC	Undef	Abort
r8				
r9				
r10				
r11				
r12				
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
spsr	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

Undef Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr
spsr

Banked out Registers

User	FIQ	IRQ	SVC	Abort
	r8			
	r9			
	r10			
	r11			
	r12			
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

User Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr

Banked out Registers

FIQ	IRQ	SVC	Undef	Abort
r8				
r9				
r10				
r11				
r12				
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
spsr	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

Abort Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr
spsr

Banked out Registers

User	FIQ	IRQ	SVC	Undef
	r8			
	r9			
	r10			
	r11			
	r12			
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
	spsr	spsr	spsr	spsr

The ARM Register Set

Current Visible Registers

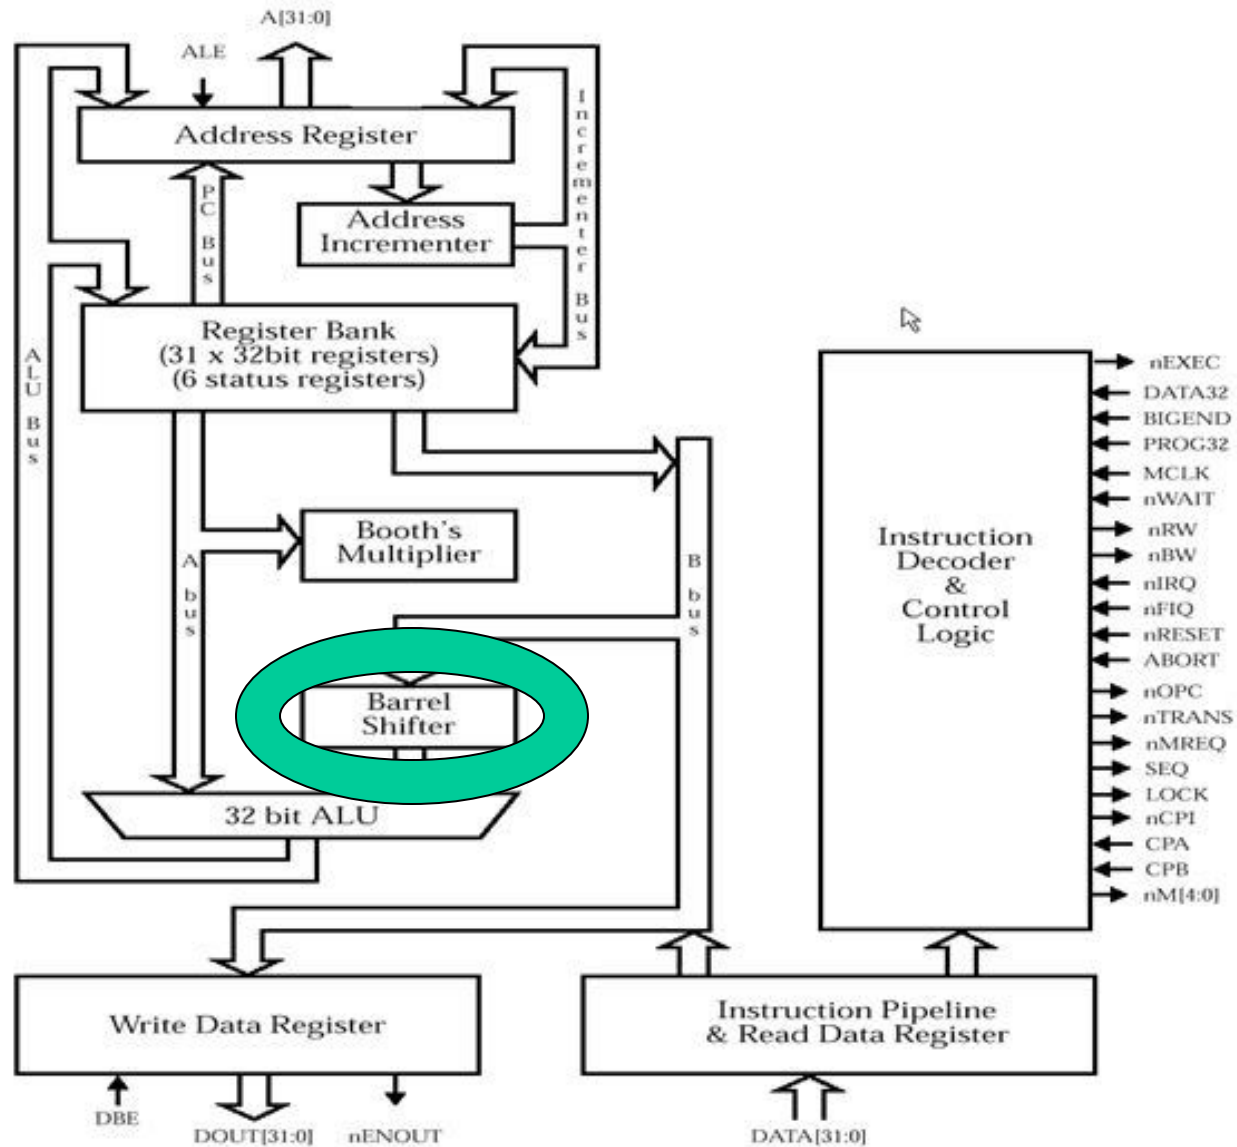
User Mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)
cpsr

Banked out Registers


FIQ	IRQ	SVC	Undef	Abort
r8				
r9				
r10				
r11				
r12				
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
spsr	spsr	spsr	spsr	spsr

ARMv7 Data Path



Other “goodies”

The ALU in ARM has a “free shift” in front of it:



Logical shift left (lsl)

Logical shift right (lsr)

Arithmetic shift right (asr)

Rotate right (ror)

Conditional execution

In LC-3, the BR instruction uses the condition code register NZP

In ARM, the condition code register is called the CPSR (Current Program Status Register). It has:

N - Negative Result from ALU

Z - Zero result from ALU

C - ALU operation had a carry-out of MSbit

V - ALU operation had a twos-complement overflow
(plus bits to handle interrupts, etc.)

In ARM, *ALL* instructions can “test” these bits

Branching...

B label or BL label (branch and link)

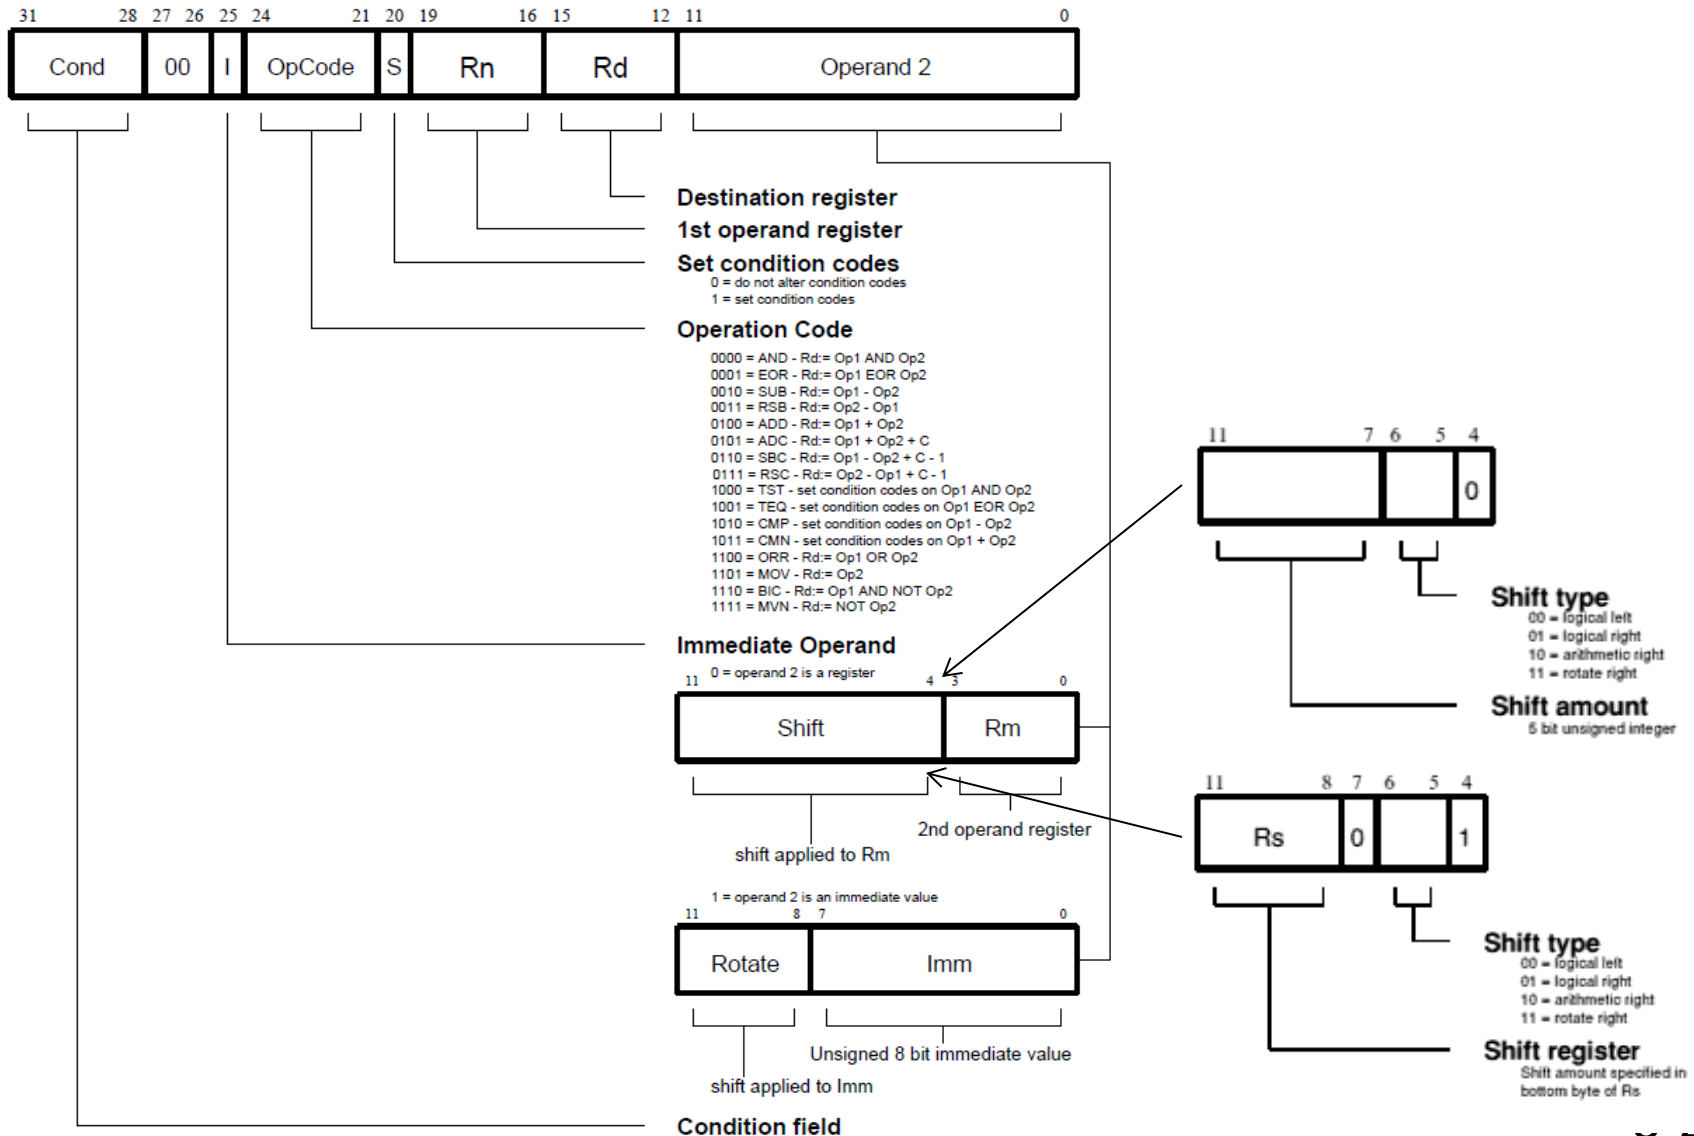
Conditional branching *Bsuffix* label...

Suffix	Flags	Meaning
EQ	Z set	Equal
NE	Z clear	Not equal
CS or HS	C set	Higher or same (unsigned \geq)
CC or LO	C clear	Lower (unsigned $<$)
MI	N set	Negative
PL	N clear	Positive or zero

more...

Suffix	Flags	Meaning
VS	V set	Overflow
VC	V clear	No overflow
HI	C set and Z clear	Higher (unsigned >)
LS	C clear or Z set	Lower or same (unsigned <=)
GE	N and V the same	Signed >=
LT	N and V differ	Signed <
GT	Z clear, N and V the same	Signed >
LE	Z set, N and V differ	Signed <=
AL	Any	Always. This suffix is normally omitted.

ARM data processing instruction format



Putting it all together: the ARM ADD instruction

ADD{S}<C> <Rd>, <Rn>, <Rm>, <type> <Rs>

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
cond				0	0	0	0	1	0	0	S	Rn				Rd				Rs				0	type	1	Rm				

```
d = UInt(Rd); n = UInt(Rn); m = UInt(Rm); s = UInt(Rs);  
setflags = (S == '1'); shift_t = DecodeRegShift(type);  
if d == 15 || n == 15 || m == 15 || s == 15 then UNPREDICTABLE;
```

Ignore “cond” for now

type = lsl, lsr, asr, or ror and applies to Rm

Examples:

ADD R1,R2,R3

r1 = r2+r3

ADDS R1,R2,R3

r1 = r2+r3 and the flags are set

ADDS R1,R2,R3 lsl R4

r1 = r2+(r3 lsl r4)

Arithmetic opcodes that work like ADD

op r1, r2, r3

SUB $r1 = r2 - r3$ **Subtract**

RSB $r1 = r3 - r2$ **Reverse subtract**

AND $r1 = r1 \& r2$ **And**

ORR $r1 = r1 | r2$ **Or**

EOR $r1 = r1 \wedge r2$ **Exclusive or**

not a comprehensive list!

ARM has about 200 instructions depending on how you count!

Some two-operand opcodes

op r1, r2

CMP set flags based on r1-r2

MOV r1, r2

note can use an immediate if needed...

CMP r1,r2

MOV r1,#0

and shift...

MOV r1,r2, lsl #2

Loads and stores

LDR*type* R1, [R2, #imm]

STR*type* R1, [R2, #imm]

What's *type*? Something LC-3 didn't have:

type is one of

- B for byte (8b)
- H for halfword (16b)
- (blank) for word (32b)
- D for doubleword (64b)

Also can add an “S” for *sign extension*

so you can have

LDRSB R1, [R2, #4] ; Loads a byte into a 32-bit register and sign extends it

All kinds of ways to get to memory...

[Rx, #imm]

[Rx, +Ry] [Rx, -Ry]

[Rx, \pm Ry, *shifto*p #imm]

also add ! to update the Ry register afterwards

LDRD R1,[R2, -R3, lsl #3]!

effective address is $R2-(R3 \ll 3)$

after the load executes: $R3 = R2-(R3 \ll 3)$

or this syntax to update the register before the load

LDRD R1,R2,[-R3] lsl #3

$R3 = R2-R3 \ll 3$ *before* the load executes

then the effective address is R3

How do you get an address in a register?

Same conundrum as in LC-3...

In LC-3

LEA R1,*label*

In ARM

LDR R1,*=label* puts the value of “label” in R1

Calls and returns

A call is a BL *label* for “branch and link”

LR (R14) gets the current PC+4

A return is via a BX *Rx*, or *usually* BX LR



Quick example

Count the number of positive values of an array of ints, A[],there are 16 values to investigate...

	LDR R0,=A	;Address of A in R0
	MOV R1,#15	;Countdown value
	MOV R2,#0	;Initialize counter
Label:	LDR R3,[R0,+R1 lsl #2]	;Get an array val
	CMP R3,#0	;Test it
	BLT Noadd	;Ignore neg vals
	ADD R2,R2,#1	;Increment count
Noadd:	SUBS R1,R1,#1	;Dec counter
	BGE Label	;More if not done

Ok some weird (cool) stuff

ARM has conditional execution for (nearly) all instructions.

Those conditions (LT, LE, GT, ...) can be appended to an opcode that allows it

So instead of this

```
CMP R3,#0
BLT Noadd
ADD R2,R2,#1
```

Noadd: ...

we could have done this

```
CMP R3,#0
ADDGE R2,R2,#1
```

Quick example

Count the number of positive values of an array, there are 16 values to investigate...

```

                LDR R0,=A
                MOV R1,#15
                MOV R2,#0
Label:          LDR R3,[R0,+R1 lsl #2]
                CMP R3,#0
                ADDGE R2,R2,#1           ;Increment if >= 0
                SUBS R1,R1,#1
                BGE Label
```

Stacks!

Often load or store multiple...

STMxx sp!, {register list} reg order doesn't matter here
! means post-update SP, xx for + or -

LDMxx sp!, {register list}

where xx is the kind of stack

FD - full (sp points to top of stack), descending (--sp to push)**

FA - full, ascending (sp++ to push)

also

ED – empty descending

EA – empty ascending

** what the ARM calling convention uses, same as LC-3

Unix calling convention for ARM

r0-r3 are the argument and scratch registers;

r0-r1 are also the return value registers

r4-r8, r10-r11 are callee-save registers

r9 special use

r11 is the frame pointer

r12 is a temporary (caller-save)

r13 the link register

r14 the stack pointer

r15 the PC

A cross compiler you can get

For Mac OSX:

<http://sourceforge.net/projects/yagarto/>

**free, pre-built binaries. Be sure to read the “read me first”
before you install**

then try

**eabi-none-arm-gcc -S prog.c
makes a prog.s assembly file**

Using GCC cross compiler...

```
int a[16];  
main() {  
    int i, j;  
  
    for (i = 15; i >= 0; i--)  
        if (a[i] >= 0) j++;  
  
    printf("%d\n", j);  
}
```

.LC0: ascii "%d\012\000"

.text

main: stmfd sp!, {fp, lr}

add fp, sp, #4

sub sp, sp, #8

mov r3, #15

str r3, [fp, #-8]

b .L2

.L4: ldr r3, .L5

ldr r2, [fp, #-8]

ldr r3, [r3, r2, lsl #2]

cmp r3, #0

blt .L3

ldr r3, [fp, #-12]

add r3, r3, #1

str r3, [fp, #-12]

.L3: ldr r3, [fp, #-8]

sub r3, r3, #1

str r3, [fp, #-8]

.L2: ldr r3, [fp, #-8]

cmp r3, #0

bge .L4

ldr r0, .L6

ldr r1, [fp, #-12]

bl printf

sub sp, fp, #4

ldmfd sp!, {fp, lr}

bx lr

.L5: .word a

.L6 .word .LC0

; save fp and lr (store multiple)

; set up frame pointer ☺

; int i, j //make space for two local vars

; for (i = 15; ...

;

; //go to test part of for loop

; if (a[i] >= 0) //load in base of 'a'

; //get i

; //get a[i]

;

; //branch less than over if part

; j++ //if part: get j

; //increment j

; //update j.

; ...i--) //reinitialize part of for

; //decrement i

; //update i

; ...i >= 0;... // test part of for loop

;

; //back to loop body if i >= 0

; printf("%d\n, j)//get the address of the format string for printf

; //get the value of j

; //branch and link (call) printf

; pop the stack frame

; restore registers

; return

; pointer to a[0]

; pointer to literal "%d\n"

Thumb

ROM is precious on an embedded device

Can “save space” in ROM using *Thumb*

Thumb is a proper subset of ARM ISA

Instructions are 16b wide, which reduces what you can do.

Here's a Thumb ADD encoding:

Encoding T1 ARMv4T, ARMv5T*, ARMv6*, ARMv7

ADDS <Rd>, <Rn>, <Rm>

Outside IT block.

ADD<c> <Rd>, <Rn>, <Rm>

Inside IT block.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	0	0	1	1	0	0	Rm			Rn			Rd		

```
d = UInt(Rd); n = UInt(Rn); m = UInt(Rm); setflags = !InITBlock();  
(shift_t, shift_n) = (SRTYPE_LSL, 0);
```

Encoding T2 ARMv6T2, ARMv7 if <Rdn> and <Rm> are both from R0-R7



ARMv4T, ARMv5T*, ARMv6*, ARMv7 otherwise

ADD<c> <Rdn>, <Rm>



If <Rdn> is the PC, must be outside or last in IT block.

15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
0	1	0	0	0	1	0	0	DN	Rm			Rdn			

```
if (DN:Rdn) == '1101' || Rm == '1101' then SEE ADD (SP plus register);  
d = UInt(DN:Rdn); n = d; m = UInt(Rm); setflags = FALSE; (shift_t, shift_n) = (SRTYPE_LSL, 0);  
if n == 15 && m == 15 then UNPREDICTABLE;  
if d == 15 && InITBlock() && !LastInITBlock() then UNPREDICTABLE;
```



ARM



Thumb

Has Push, Pop

ARM

```
main: stmfd sp!, {fp, lr}
      add  fp, sp, #4
      sub  sp, sp, #8
      mov  r3, #15
      str  r3, [fp, #-8]
      b    .L2
.L4:   ldr  r3, .L5
      ldr  r2, [fp, #-8]
      ldr  r3, [r3, r2, lsl #2]
      cmp  r3, #0
      blt  .L3
      ldr  r3, [fp, #-12]
      add  r3, r3, #1
      str  r3, [fp, #-12]
.L3:   ldr  r3, [fp, #-8]
      sub  r3, r3, #1
      str  r3, [fp, #-8]
.L2:   ldr  r3, [fp, #-8]
      cmp  r3, #0
      bge  .L4
      ldr  r0, .L6
      ldr  r1, [fp, #-12]
      bl   printf
      mov  r0, r3
      sub  sp, fp, #4
      ldmfd sp!, {fp, lr}
      bx   lr
```

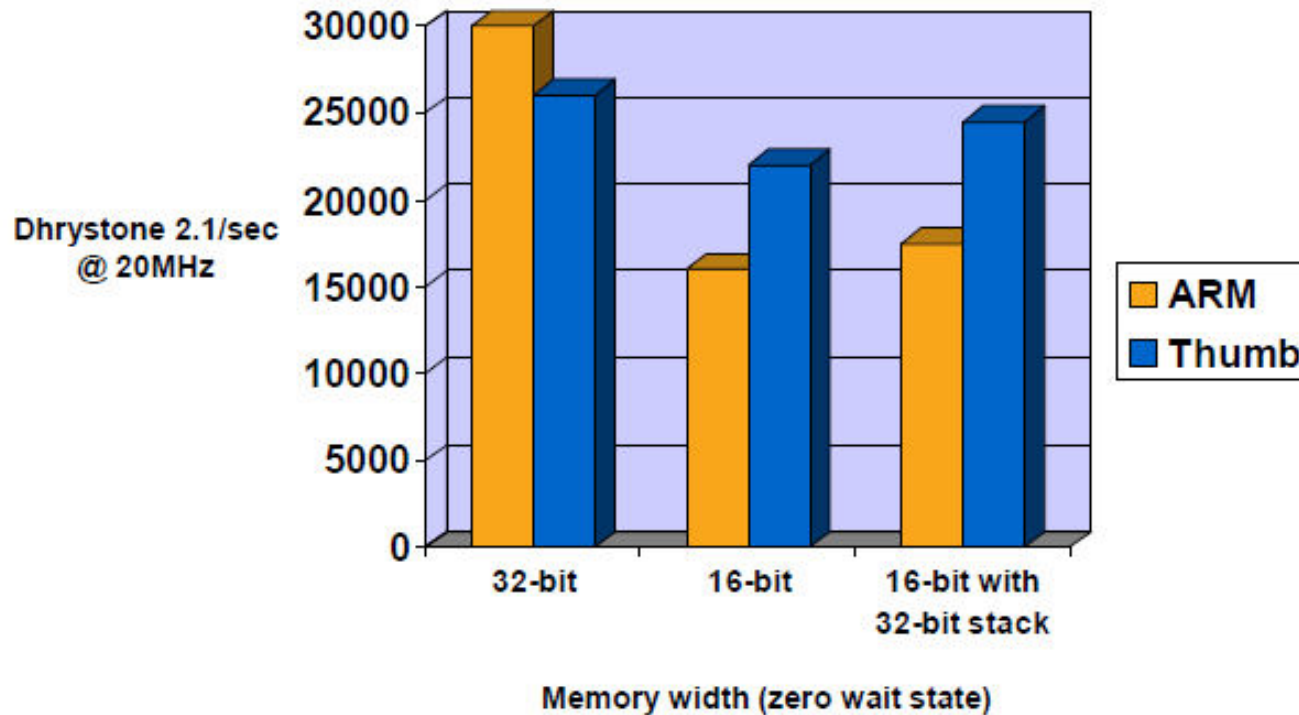
27 instructions x 4B each
total = 108 Bytes

Thumb

```
main:  push {r7, lr}
      sub  sp, sp, #8
      add  r7, sp, #0
      mov  r3, #15
      str  r3, [r7, #4]
      b    .L2
.L4:    ldr  r3, .L5
      ldr  r2, [r7, #4]
      lsl  r2, r2, #2
      ldr  r3, [r2, r3]
      cmp  r3, #0
      blt  .L3
      ldr  r3, [r7]
      add  r3, r3, #1
      str  r3, [r7]
.L3:    ldr  r3, [r7, #4]
      sub  r3, r3, #1
      str  r3, [r7, #4]
.L2:    ldr  r3, [r7, #4]
      cmp  r3, #0
      bge  .L4
      ldr  r2, .L6
      ldr  r3, [r7]
      mov  r0, r2
      mov  r1, r3
      bl   printf
      mov  r0, r3
      mov  sp, r7
      add  sp, sp, #8
      pop  {r7}
      pop  {r1}
      bx   r1
```

32 instructions x 2B each
total = 64 Bytes

ARM and Thumb Performance



Floating Point (!!!)

Uses a separate register file

Has either 32 32b registers called S0 to S31
or 16 64b registers called D0 to D15

To get values from “regular” registers in/out of these, you
can use VMOV

VMOV.*size* S3, R1

where *size* is .8 or .16 or .32

VADD.F32 Sd, Sn, Sm also VMUL, VDIV, VSUB, VSQRT...

VADD.F64 Dd, Dn, Dm

Smart C Programming

First 4 parameters passed in registers

In loops always make the loop variable count down, if possible

Size of local variables

```
int wordsize(int a)
```

```
{  
    a=a+1;  
    return a;  
}
```

```
int halfsize(short b)
```

```
{  
    b=b+1;  
    return b;  
}
```

```
int bytesize(char c)
```

```
{  
    c=c+1;  
    return c;  
}
```

```
wordsize
```

```
ADD r0,r0,#1
```

```
BX lr
```

```
halfsize
```

```
ADD r0,r0,#1
```

```
MOV r0,r0,LSL #16
```

```
MOV r0,r0,ASR #16
```

```
BX lr
```

```
bytesize
```

```
ADD r0,r0,#1
```

```
AND r0,r0,#0xFF
```

```
BX lr
```

Resources

<http://www.arm.com/support/university/>